

Alpenhorn

Preliminaries

1. Start the DB and enter the first container

```
docker-compose up -d db
docker-compose run host_1 bash -l
```

2. Initialize Alpenhorn database

```
alpenhorn init
alpenhorn create_group group_1
alpenhorn create_node --auto_import=yes node_1 /data host_1 group_1
alpenhorn mount --user root --address host_1 node_1
```

Importing files

3. Start the service

```
docker-compose up host_1
```

Note files being imported. (Could also do `docker-compose logs host_1 | head -30.`)

4. Client

```
docker-compose exec host_1 bash -l
alpenhorn status
```

Two files on host_1

Start more nodes

5. Create storage

```
alpenhorn create_group group_2
alpenhorn create_node --auto_import=yes node_2 /data host_2 group_2
alpenhorn mount --user root --address host_2 --hostname host_2 node_2
```

```
alpenhorn create_group group_3
alpenhorn create_node --auto_import=yes node_3 /data host_3 group_3
alpenhorn mount --user root --address host_3 --hostname host_3 node_3
```

6. Start the service

```
docker-compose up -d host_1
docker-compose up host_2
docker-compose up host_3
```

Note no files are being imported.

Syncing files

7. Client

```
docker-compose exec host_2 bash -l
alpenhorn verify
find /data
```

No files on host_2

8. Start a sync

```
alpenhorn sync node_1 group_2 --show_files
```

Note the server copying

```
find /data
```

The files are here!

9. Sync an explicit acquisition

```
docker-compose exec host_3 bash -l
find /data
alpenhorn sync node_1 group_3 --acq=12345678T000000Z_inst_zab --show_files
```

Note the server copying

10. Client

```
alpenhorn status
alpenhorn verify node_2
alpenhorn verify node_3
```

Two files on host_2, one file on host_3.

Add a new file

11. create a new file

```
docker-compose exec host_2 bash -l
echo foo bar > /data/12345678T000000Z_inst_zab/jim.out
```

12. see the new file get synced

```
alpenhorn sync --acq 12345678T000000Z_inst_zab node_2 group_3 --show_files
alpenhorn status
docker-compose up host_3
```

Highlight “transferring file ... jim.out”

13. See how many files are present

```
alpenhorn status
```

Dealing with corruption

14. Modify a copy of a file:

```
alpenhorn status
echo bla >> /data/12345678T000000Z_inst_zab/jim.out
alpenhorn status
```

Notice in the output of the service:

```
> INFO >> Checking file "/data/12345678T000000Z_inst_zab/jim.out" on node "node_2".
> ERROR >> File is corrupted!
> INFO >> Updating file copy status [id=6].
```

15. Repair with a sync from a known good copy:

```
alpenhorn sync --acq 12345678T000000Z_inst_zab node_1 group_2 --show_files
alpenhorn status
```

Cleaning

16. Remove unneeded files with clean:

```
alpenhorn clean --now node_2
```

Observe only two are removed: “Too few backups to delete
2017/03/21/acq_xy1_45678901T000000Z_inst_zab/acq_data/x_123_1_data/raw/acq_123_1.zxc”.
This is because this acquisition was never synced to node_3

Transport disks

17. Create storage

```
alpenhorn create_group transport
alpenhorn create_node --storage_type=T t1 /mnt/t1 host_1 transport
```

18. Make it available

```
alpenhorn mount --address=host_1 t1
```

19. Copy to transport for a target destination:

```
alpenhorn sync node_1 transport --target group_3 --show_files
```

Note files being synced

```
alpenhorn status
find /mnt/t1
alpenhorn unmount t1
```

20. On the other side (host_3)...

```
alpenhorn mount --address=host_3 t1
alpenhorn sync t1 group_3 --show_files
```

Note files being synced

```
alpenhorn status
alpenhorn clean --now t1
alpenhorn status
find /mnt/t1
```